

# The Entity-Relationship Model

## Chapter 2

**Instructor: Vladimir Zadorozhny**

[vladimir@sis.pitt.edu](mailto:vladimir@sis.pitt.edu)

Information Science Program

School of Information Sciences,

University of Pittsburgh

## Database: a Set of Relations (Tables)

| customer_id | customer_name | customer_street | customer_city |
|-------------|---------------|-----------------|---------------|
| 192-83-7465 | Johnson       | 12 Alma St.     | Palo Alto     |
| 677-89-9011 | Hayes         | 3 Main St.      | Harrison      |
| 182-73-6091 | Turner        | 123 Putnam Ave. | Stamford      |
| 321-12-3123 | Jones         | 100 Main St.    | Harrison      |
| 336-66-9999 | Lindsay       | 175 Park Ave.   | Pittsfield    |
| 019-28-3746 | Smith         | 72 North St.    | Rye           |

(a) The customer table

| account_number | balance |
|----------------|---------|
| A-101          | 500     |
| A-215          | 700     |
| A-102          | 400     |
| A-305          | 350     |
| A-201          | 900     |
| A-217          | 750     |
| A-222          | 700     |

(b) The account table

| customer_id | account_number |
|-------------|----------------|
| 192-83-7465 | A-101          |
| 192-83-7465 | A-201          |
| 019-28-3746 | A-215          |
| 677-89-9011 | A-102          |
| 182-73-6091 | A-305          |
| 321-12-3123 | A-217          |
| 336-66-9999 | A-222          |
| 019-28-3746 | A-201          |

(c) The depositor table

- Find the name of the customer with customer-id 192-83-7465  

```
select customer.customer_name
from customer
where customer.customer_id = '192-83-7465'
```

## Database Design


The process of designing the general structure of the database:

- ❖ Requires that we find a “good” collection of relation schemas.
  - Business decision – What attributes should we record in the database?
  - IS decision – What relation schemas should we have and how should the attributes be distributed among the various relation schemas?
- ❖ Deciding on the physical layout of the database

## Conceptual Database Design


- ❖ Conceptual design: (*ER Model is used at this stage.*)
  - What are the *entities* and *relationships* in the enterprise?
  - What information about these entities and relationships should we store in the database?
  - What are the *integrity constraints* or *business rules* that hold?
  - A database ‘schema’ in the ER Model can be represented pictorially (*ER diagrams*).
  - Can map an ER diagram into a relational schema.

## ER Model Basics



- ❖ **Entity:** Real-world object distinguishable from other objects. An entity is described (in DB) using a set of **attributes**.
- ❖ **Entity Set:** A collection of similar entities. E.g., all employees.
  - All entities in an entity set have the same set of attributes. (Until we consider ISA hierarchies, anyway!)
  - Each entity set has a **key**.
  - Each attribute has a **domain**.

## ER Model Basics (Contd.)





- ❖ **Relationship:** Association among two or more entities. E.g., Attishoo works in Pharmacy department.
- ❖ **Relationship Set:** Collection of similar relationships.
  - An n-ary relationship set R relates n entity sets E1 ... En; each relationship in R involves entities e1 in E1, ..., en in En
    - Same entity set could participate in different relationship sets, or in different “roles” in same set.

## Key Constraints


### ❖ Consider Works\_In:

An employee can work in many departments; a dept can have many employees.


- ❖ In contrast, each dept has at most one manager, according to the key constraint on Manages.




1-to-1



1-to Many



Many-to-1




Many-to-Many

## Participation Constraints


### ❖ Does every department have a manager?

- If so, this is a participation constraint: the participation of Departments in Manages is said to be *total* (vs. *partial*).
  - Every *did* value in Departments table must appear in a row of the Manages table (with a non-null *ssn* value!)



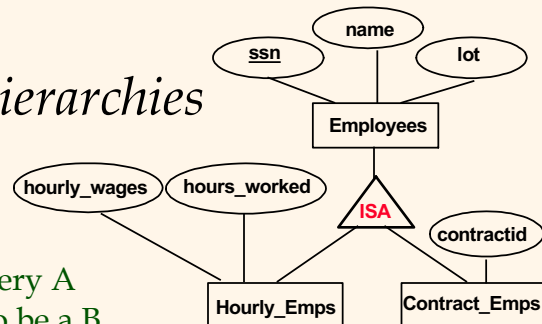
## Weak Entities

- ❖ A **weak entity** can be identified uniquely only by considering the primary key of another (*owner*) entity.
  - Owner entity set and weak entity set must participate in a one-to-many relationship set (one owner, many weak entities).
  - Weak entity set must have total participation in this **identifying** relationship set.



## ISA ('is a') Hierarchies

- ❖ As in C++, or other PLs, attributes are inherited.
- ❖ If we declare A **ISA** B, every A entity is also considered to be a B entity.



- ❖ **Overlap constraints:** Can Joe be an Hourly\_Emps as well as a Contract\_Emps entity? (*Allowed/disallowed*)
- ❖ **Covering constraints:** Does every Employees entity also have to be an Hourly\_Emps or a Contract\_Emps entity? (*Yes/no*)
- ❖ Reasons for using ISA:
  - To add descriptive attributes specific to a subclass.
  - To identify entities that participate in a relationship.

## Conceptual Design Using the ER Model

### ❖ Design choices:

- Should a concept be modeled as an entity or an attribute?
- Should a concept be modeled as an entity or a relationship?

### ❖ Constraints in the ER Model:

- A lot of data semantics can (and should) be captured.
- But some constraints cannot be captured in ER diagrams.

## Summary of Conceptual Design

- ❖ *Conceptual design follows requirements analysis,*
  - Yields a high-level description of data to be stored
- ❖ ER model popular for conceptual design
  - Constructs are expressive, close to the way people think about their applications.
- ❖ Basic constructs: *entities, relationships, and attributes* (of entities and relationships).
- ❖ Some additional constructs: *weak entities, ISA hierarchies.*
- ❖ Note: There are many variations on ER model.

## Summary of ER (Contd.)


- ❖ Several kinds of integrity constraints can be expressed in the ER model: *key constraints*, *participation constraints*, and *overlap/covering constraints* for ISA hierarchies. Some *foreign key constraints* are also implicit in the definition of a relationship set.
  - Some constraints (notably, *functional dependencies*) cannot be expressed in the ER model.
  - Constraints play an important role in determining the best database design for an enterprise.

## Summary of ER (Contd.)

- ❖ ER design is *subjective*. There are often many ways to model a given scenario! Analyzing alternatives can be tricky, especially for a large enterprise. Common choices include:
  - Entity vs. attribute, entity vs. relationship, whether or not to use ISA hierarchies.
- ❖ Ensuring good database design: resulting relational schema should be analyzed and refined further. FD information and normalization techniques are especially useful.


## Logical DB Design: ER to Relational

### ❖ Entity sets to tables:



```
CREATE TABLE Employees
(ssn CHAR(11),
name CHAR(20),
lot INTEGER,
PRIMARY KEY (ssn))
```

## Relationship Sets to Tables



### ❖ In translating a relationship set to a relation, attributes of the relation must include:



- Keys for each participating entity set (as foreign keys).
  - This set of attributes forms a *superkey* for the relation.
- All descriptive attributes.

```
CREATE TABLE Works_In(
ssn CHAR(11),
did INTEGER,
since DATE,
PRIMARY KEY (ssn, did),
FOREIGN KEY (ssn)
REFERENCES Employees,
FOREIGN KEY (did)
REFERENCES Departments)
```




## Review: Key Constraints


- ❖ Each dept has at most one manager, according to the key constraint on Manages.




1-to-1



1-to Many



Many-to-1



Many-to-Many

*Translation to relational model?*

## Translating ER Diagrams with Key Constraints

- ❖ Map relationship to a table:
  - Note that **did** is the key now!
  - Separate tables for Employees and Departments.
- ❖ Since each department has a unique manager, we could instead combine Manages and Departments.


```
CREATE TABLE Manages(
  ssn CHAR(11),
  did INTEGER,
  since DATE,
  PRIMARY KEY (did),
  FOREIGN KEY (ssn) REFERENCES Employees,
  FOREIGN KEY (did) REFERENCES Departments)
```

```
CREATE TABLE Dept_Mgr(
  did INTEGER,
  dname CHAR(20),
  budget REAL,
  ssn CHAR(11),
  since DATE,
  PRIMARY KEY (did),
  FOREIGN KEY (ssn) REFERENCES Employees)
```

## Review: Participation Constraints

### ❖ Does every department have a manager?

- If so, this is a participation constraint: the participation of Departments in Manages is said to be *total* (vs. *partial*).
  - Every *did* value in Departments table must appear in a row of the Manages table (with a non-null *ssn* value!)




## Participation Constraints in SQL

- ### ❖ We can capture participation constraints involving one entity set in a binary relationship, but little else (without resorting to CHECK constraints).

```
CREATE TABLE Dept_Mgr(  
  did INTEGER,  
  dname CHAR(20),  
  budget REAL,  
  ssn CHAR(11) NOT NULL,  
  since DATE,  
  PRIMARY KEY (did),  
  FOREIGN KEY (ssn) REFERENCES Employees,  
  ON DELETE NO ACTION)
```

## Review: Weak Entities

- ❖ A *weak entity* can be identified uniquely only by considering the primary key of another (*owner*) entity.
  - Owner entity set and weak entity set must participate in a one-to-many relationship set (1 owner, many weak entities).
  - Weak entity set must have total participation in this *identifying* relationship set.

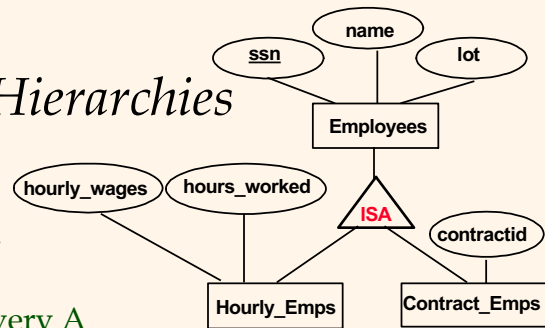


## Translating Weak Entity Sets

- ❖ Weak entity set and identifying relationship set are translated into a single table.
  - When the owner entity is deleted, all owned weak entities must also be deleted.

```
CREATE TABLE Dep_Policy (  
  pname CHAR(20),  
  age INTEGER,  
  cost REAL,  
  ssn CHAR(11) NOT NULL,  
  PRIMARY KEY (pname, ssn),  
  FOREIGN KEY (ssn) REFERENCES Employees,  
  ON DELETE CASCADE)
```

## Review: ISA Hierarchies



- ❖ As in C++, or other PLs, attributes are inherited.
- ❖ If we declare A **ISA** B, every A entity is also considered to be a B entity.
- ❖ **Overlap constraints:** Can Joe be an Hourly\_Emps as well as a Contract\_Emps entity? (*Allowed/disallowed*)
- ❖ **Covering constraints:** Does every Employees entity also have to be an Hourly\_Emps or a Contract\_Emps entity? (*Yes/no*)

## Translating ISA Hierarchies to Relations

- ❖ **General approach:**
  - 3 relations: Employees, Hourly\_Emps and Contract\_Emps.
    - Hourly\_Emps: Every employee is recorded in Employees. For hourly emps, extra info recorded in Hourly\_Emps (*hourly\_wages, hours\_worked, ssn*); must delete Hourly\_Emps tuple if referenced Employees tuple is deleted).
    - Queries involving all employees easy, those involving just Hourly\_Emps require a join to get some attributes.
- ❖ **Alternative: Just Hourly\_Emps and Contract\_Emps.**
  - Hourly\_Emps: *ssn, name, lot, hourly\_wages, hours\_worked*.
  - Each employee must be in one of these two subclasses.